

## Socket

CIS 340 System Programming  
Haodong Wang

## New IPC Ideas

- IPC: inter-process communication
- Until now, we did the communication through:
  - File system: file descriptors, read()/write() and pipes
  - Total reliable byte stream between producer and consumer
- New ideas:
  - We want to create a generalization of the pipe construct for network-based I/O
  - That means we still want file descriptors and read()/write() calls to work
  - We need to take some extra features into consideration:
    - Network Protocol Stacks
    - Network Naming Conventions
    - Requirements of Protocol-Specific Message Passing
  - The BSD and UNIX solution is the socket() call. Most concisely, it can be described as a communication endpoint.
  - The call returns a file descriptor.
  - int socket(int domain, int type, int protocol);

4/2/2015

CIS 340 Systems Programming

2

## Communication Domain

- This basically specifies a protocol stack.
- Some systems contain a richer set of communication domains than others
  - AF\_UNIX or AF\_LOCAL: The UNIX IPC domain, local to a single machine
  - AF\_INET: The Internet domain, global in scope
  - AF\_INET6: The Internet domain, using IPv6
- Once a domain is specified, we know how to associate a name with the socket
- As well as knowing the semantics of supported IPC mechanisms

## Unix Domain Sockets

- Let's start with the simpler (but less interesting) case of the AF\_UNIX communication domain
- The header file <sys/un.h> defines addresses

```
#define UNIX_PATH_MAX 108

struct sockaddr_un {
    unsigned short sun_family; /* AF_UNIX */
    char sun_path[UNIX_PATH_MAX]; /* Pathname */
};
```

- Some examples of Unix domain sockets can be found under the /dev directory

4/2/2015

CIS 340 Systems Programming

3

4/2/2015

CIS 340 Systems Programming

4

## Types of Sockets in Unix Domain

- We'll be looking at two types of sockets available in Unix
  - SOCK\_DGRAM provides datagram communication semantics
  - Only promises best-effort delivery
  - Unix may discard datagrams in times of buffer congestion
  - Connectionless!
- 2<sup>nd</sup> type
  - SOCK\_STREAM implements virtual circuit communication semantics
  - Reliable FIFO point-to-point communications
  - Appears as a byte stream to applications
  - This is actually how some later UNIX systems implement pipes!
- Socket type should be chosen according to the needs of the application, and should be programmed in accordance with well-specified delivery semantics of chosen type.

4/2/2015

CIS 340 Systems Programming

5

## Operations on Sockets

- Binding a name to a socket:

```
int bind(int sockfd, struct sockaddr *my_addr, int addrlen);
```

- Sending datagram on a socket (asynchronous):

```
int sendto(int s, const void *msg, int len, unsigned int flags,
           const struct sockaddr *to, int tolen);
```

- Receiving datagram from a socket (synchronous, blocking):

```
int recvfrom(int s, void *buf, int len, unsigned int flags,
             struct sockaddr *from, int *fromlen);
```

4/2/2015

CIS 340 Systems Programming

6

## Server I

```
#include <errno.h>
#include <strings.h>
#include <stdio.h>
#include <unistd.h>
#include <sys/socket.h>
#include <sys/un.h>

main() {
    short p_len;
    int socket_fd, cc, h_len, fsize, namelen;
    void printsun();
    struct sockaddr_un s_un, from;
    size_t addrlen;

    struct {
        char head;
        u_long body;
        char tail;
    } msg;

    socket_fd = socket (AF_UNIX, SOCK_DGRAM, 0);
    s_un.sun_family = AF_UNIX;
```

4/2/2015

CIS 340 Systems Programming

7

## Server II

```
strcpy(s_un.sun_path, "udgram");
addrlen = sizeof(s_un.sun_family) + sizeof(s_un.sun_path);
unlink("udgram");

bind(socket_fd, (struct sockaddr *)&s_un, addrlen);

for(;;) {
    fsize = sizeof(from);
    cc = recvfrom(socket_fd, &msg, sizeof(msg), 0,
                 (struct sockaddr *)&from, &fsize);
    printsun(&from, "unix_idgram:", "Packet,from");
    printf("Got data::%c%d%c\n", msg.head, msg.body, msg.tail);
    fflush(stdout);
}

void printsun(Sun, s1, s2)
struct sockaddr_un *Sun; char *s1, *s2;
{
    printf("%s,%s\n", s1, s2);
    printf("uuuuuu family,<id>,<addr,<%s>\n",
           Sun->sun_family, Sun->sun_path);
}
```

4/2/2015

CIS 340 Systems Programming

8

## Client I

```
#include <errno.h>
#include <strings.h>
#include <stdio.h>
#include <unistd.h>
#include <sys/socket.h>
#include <sys/un.h>

main() {
    int socket_fd, cc;
    //long getpid();
    struct sockaddr_un dest;

    struct {
        char head;
        u_long body;
        char tail;
    } msgbuf;

    socket_fd = socket (AF_UNIX, SOCK_DGRAM, 0);
    dest.sun_family = AF_UNIX;
    strcpy(dest.sun_path, "udgram");

    msgbuf.head = '<';
```

4/2/2015

CIS 340 Systems Programming

9

## Client II

```
msgbuf.body = (u_long) getpid();
msgbuf.tail = '>';

cc = sendto(socket_fd, &msgbuf, sizeof(msgbuf), 0,
            (struct sockaddr *)&dest, sizeof(dest));
}
```

4/2/2015

CIS 340 Systems Programming

10

## Sockets and the Internet (IPv4)

- AF\_INET communication domain
- SOCK\_DGRAM - Same as before! UDP/IP
- SOCK\_STREAM - Same as before! TCP/IP
- We need a way to associate names with sockets to be able to do network I/O through a socket file descriptor

4/2/2015

CIS 340 Systems Programming

11

## Sockets and the Internet (IPv4)

- The header file <netinet/in.h> defines a 32-bit for an Internet host.
- This actually identifies a specific network interface on a specific system on the Internet.
- It's represented by a 32-bit unsigned number

```
struct in_addr {
    __u32 s_addr;
}
```

- The addresses are usually represented by dotted decimal notation.

4/2/2015

CIS 340 Systems Programming

12

## Representing the Address in C

- In header file <netinet/in.h>
 

```
#define _SOCK_SIZE_ 16 /* sizeof(struct sockaddr) */

struct sockaddr_in {
    short int sin_family; /* Address family */
    unsigned short int sin_port; /* Port number */
    struct in_addr sin_addr; /* Internet address */
    /* Pad to size of "struct sockaddr" */
    unsigned char _pad[_SOCK_SIZE_ - sizeof(short int) -
        sizeof(unsigned short int) - sizeof(struct in_addr)];
};
```
- Declare and/or allocate instance of struct sockaddr\_in whenever you need to specify a full address on the Internet
- A port is an Internet communication endpoint associated with an application. (host,port) defines an Internet address.
- Ports in range [0,1023] reserved for root; others available to ordinary users. (See RFC 1700)

4/2/2015

CIS 340 Systems Programming

13

## Usual Ports for Services

- FTP uses 20 and 21
- SSH uses 22
- Telnet uses 23
- HTTP uses 80, commonly
- HTTPS uses 443
- Check /etc/services to see what "well-known" ports are

4/2/2015

CIS 340 Systems Programming

14

## Translating Host Names into IP Address(es)

- Library function to map symbolic host name into IP address(es):
 

```
#include <netdb.h>

struct hostent *gethostbyname(const char *name);

void herror(const char *s);
```
- The hostent data structure:
 

```
struct hostent {
    char *h_name; /* official name of host */
    char **h_aliases; /* alias list */
    int h_addrtype; /* host address type */
    int h_length; /* length of address */
    char **h_addr_list; /* list of addresses */
};
#define h_addr h_addr_list[0] /* for backward compatibility */
```

4/2/2015

CIS 340 Systems Programming

15

## Translating Host Names into IP Address(es)

- If we have a dotted decimal string and we want to convert it into an address we can use, the above function is useful.
- I Also see man inet for more functions!

4/2/2015

CIS 340 Systems Programming

16

## getaddrs.c Example

```
#include <netdb.h>
#include <stdio.h>
#include <stdlib.h>
#include <sys/socket.h>
#include <arpa/inet.h>
#include <netinet/in.h>
main(argc, argv)
int argc; char **argv;
{
    struct hostent *entry; char **next;
    struct in_addr address; **addrptr;
    entry = gethostbyname(argv[1]);
    if (!entry) { perror("lookup_error"); exit(1); }
    printf("Official_name->,%s\n", entry->h_name);
    if (entry->h_aliases[0]) {
        printf("Aliases->\n");
        for (next = entry->h_aliases; *next; next++)
            printf("%s\n", *next);
    }
    printf("IP_Addresses:\n");
    for (addrptr=(struct in_addr *) entry->h_addr_list;
        *addrptr; addrptr++)
        printf("%s\n", inet_ntoa(*addrptr));
}
```

4/2/2015

CIS 340 Systems Programming

17

## Get Symbolic Name from IP Address

- There's also an inverse function (we know IP address, want symbolic name)

```
#include <netdb.h>

struct hostent *gethostbyaddr(const char *addr,
                              int len, int type);
```

4/2/2015

CIS 340 Systems Programming

18

## gethost.c Example


```
#include <netdb.h>
#include <stdio.h>
#include <stdlib.h>
#include <sys/socket.h>
#include <arpa/inet.h>
#include <netinet/in.h>
main(argc, argv)
int argc; char **argv;
{
    struct hostent *entry; char **next;
    struct in_addr address; **addrptr;
    inet_aton(argv[1], &address);
    entry = gethostbyaddr((char *)&address, sizeof(address),
        AF_INET);
    if (!entry) { perror("lookup_error"); exit(1); }
    printf("Official_name->,%s\n", entry->h_name);
    if (entry->h_aliases[0]) {
        printf("Aliases->\n");
        for (next = entry->h_aliases; *next; next++)
            printf("%s\n", *next);
    }
    printf("IP_Addresses:\n");
    for (addrptr=(struct in_addr *) entry->h_addr_list;
        *addrptr; addrptr++)
        printf("%s\n", inet_ntoa(*addrptr));
}
```

4/2/2015

CIS 340 Systems Programming

19

## Communication Architecture



4/2/2015

CIS 340 Systems Programming

20

## recv\_upd.c: UDP/IP Server I

```
#include <netdb.h>
#include <stdio.h>
#include <stdlib.h>
#include <sys/socket.h>
#include <arpa/inet.h>
#include <netinet/in.h>
#include <strings.h>
void printsin(s_in, s1, s2)
{
    struct sockaddr_in *s_in; char *s1, *s2;
    {
        printf ("Program: %s\n%s", s1, s2);
        printf ("%d,%d\n", s_in->sin_addr.s_addr, s_in->sin_port);
    }
}

main()
{
    short p_len;
    int socket_fd, cc, h_len, fsize, namelen;
    struct sockaddr_in s_in, from;
    struct { char head; u_long body; char tail; } msg;

    socket_fd = socket (AF_INET, SOCK_DGRAM, 0);
    /* You must do this just in case */
```

4/2/2015

CIS 340 Systems Programming

21

## recv\_upd.c: UDP/IP Server II

```
bzero((char *) &s_in, sizeof(s_in));

s_in.sin_family = (short)AF_INET;
s_in.sin_addr.s_addr = htonl(INADDR_ANY); /* WILDCARD */
s_in.sin_port = htons((u_short)0x3333);
printf (&s_in, "RECV_UDP", "Local socket is:");
fflush(stdout);
bind(socket_fd, (struct sockaddr *)&s_in, sizeof(s_in));
for(;;) {
    fsize = sizeof(from);
    cc = recvfrom(socket_fd, &msg, sizeof(msg), 0,
        (struct sockaddr *)&from, &fsize);
    printf (&from, "recv_udp: ", "Packet from:");
    printf ("Got data::%c%d%c\n", msg.head,
        ntohs(msg.body), msg.tail);
    fflush(stdout);
}
```

4/2/2015

CIS 340 Systems Programming

22

## send\_upd.c: UDP/IP Client I

```
#include <netdb.h>
#include <stdio.h>
#include <stdlib.h>
#include <sys/socket.h>
#include <arpa/inet.h>
#include <netinet/in.h>
#include <strings.h>
main(argc, argv)
int argc; char **argv;
{
    int socket_fd;
    struct sockaddr_in dest;
    struct hostent *hostptr;
    struct { char head; u_long body; char tail; } msgbuf;
    socket_fd = socket (AF_INET, SOCK_DGRAM, 0);
    /* You must do this just in case */
    bzero((char *) &dest, sizeof(dest));

    hostptr = gethostbyname(argv[1]);
    dest.sin_family = (short) AF_INET;
    bcopy(hostptr->h_addr, (char *)&dest.sin_addr,
        hostptr->h_length);
    dest.sin_port = htons((u_short)0x3333);
```

4/2/2015

CIS 340 Systems Programming

23

## send\_upd.c: UDP/IP Client II

```
msgbuf.head = '<';
msgbuf.body = htonl(getpid()); /* IMPORTANT! */
msgbuf.tail = '>';
sendto(socket_fd, &msgbuf, sizeof(msgbuf), 0,
    (struct sockaddr *)&dest, sizeof(dest));
}
```

4/2/2015

CIS 340 Systems Programming

24

## Similarities and Differences

- Note that there are striking similarities between Unix datagram programs and Internet datagram programs
- We do need to do extra work for Internet programs
- Socket creation parameters are trivially different
- Naming conventions are significantly different
- The underlying implementation is completely different! But hidden from the programmers
- Practical note: You can always test and develop network programs on "localhost" (127.0.0.1). The implementation should be smart enough to NOT send the packets over the network (instead just pass it from output buffer to input buffer)